

In []:

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
```

In [4]:

```
df=pd.read_excel("Company_Data.csv.xlsx")
```

In [6]:

```
df.head()
```

Out[6]:

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urb
0	9.50	138.0	73.0	11.0	276.0	120.0	Bad	42.0	17.0	Y
1	11.22	111.0	48.0	16.0	260.0	83.0	Good	65.0	10.0	Y
2	10.06	113.0	35.0	10.0	269.0	80.0	Medium	59.0	12.0	Y
3	7.40	117.0	100.0	4.0	466.0	97.0	Medium	55.0	14.0	Y
4	4.15	141.0	64.0	3.0	340.0	128.0	Bad	38.0	13.0	Y

In [23]:

```
df.drop(['ShelveLoc','Urban','US'], axis=1 , inplace=True)
```

In [24]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sales           400 non-null    float64
1   CompPrice       400 non-null    float64
2   Income          400 non-null    float64
3   Advertising     400 non-null    float64
4   Population      400 non-null    float64
5   Price           400 non-null    float64
6   Age             400 non-null    float64
7   Education       400 non-null    float64
dtypes: float64(8)
memory usage: 25.1 KB
```

In [25]:

```
df.isna().sum()
```

Out[25]:

```
Sales          0
CompPrice      0
Income         0
Advertising    0
Population     0
Price          0
Age           0
Education      0
dtype: int64
```

In [36]:

```
x=df.drop(columns=["Education"])
y=df["Education"]
```

In [37]:

```
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=0)
```

In [38]:

```

#Bagging Meta Estimator Classifier
from sklearn.ensemble import BaggingClassifier
#making the bagging classifier with 100 decision trees
model=BaggingClassifier(n_estimators=100)
#fitting data to bagging model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))

```


```

accuracy is : 0.1375
              precision    recall  f1-score   support

    10.0         0.36         0.31         0.33         13
    11.0         0.14         0.12         0.13          8
    12.0         0.09         0.09         0.09         11
    13.0         0.14         0.11         0.12          9
    14.0         0.18         0.33         0.24          6
    15.0         0.00         0.00         0.00          8
    16.0         0.11         0.11         0.11          9
    17.0         0.06         0.33         0.11          3
    18.0         0.00         0.00         0.00         13

 accuracy                   0.14         80
macro avg                   0.12         0.16         0.13         80
weighted avg                0.13         0.14         0.13         80

```



Implementing The Bagging Classifier with Logistic Model

In [39]:

```

#Bagging Meta Estimator Classifier
from sklearn.ensemble import BaggingClassifier
from sklearn.linear_model import LogisticRegression
#making the bagging classifier with 100 Logistic Regression models
model=BaggingClassifier(base_estimator=LogisticRegression(),n_estimators=100)
#fitting data to bagging model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))

```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model_logistic.
 py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown i
 n:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.  

py:763: ConvergenceWarning: lbfgs failed to converge (status=1):  

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown i
 n:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Random Forest Classifier

In [40]:

```

from sklearn.ensemble import RandomForestClassifier
model=RandomForestClassifier(n_estimators=100)
#fitting data to Random Forest model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))

```


```

accuracy is : 0.1875
              precision    recall  f1-score   support

    10.0         0.33         0.31         0.32         13
    11.0         0.33         0.38         0.35          8
    12.0         0.08         0.09         0.09         11
    13.0         0.12         0.11         0.12          9
    14.0         0.18         0.33         0.24          6
    15.0         0.00         0.00         0.00          8
    16.0         0.22         0.22         0.22          9
    17.0         0.13         0.67         0.22          3
    18.0         0.00         0.00         0.00         13

 accuracy          0.19         0.19         0.19         80
macro avg          0.16         0.23         0.17         80
weighted avg       0.16         0.19         0.16         80

```



Tuning HyperParameters of Random Forest

In [42]:

```
#model
model=RandomForestClassifier()
n_estimators =[10,50,100,1000]
criterion =["gini", "entropy"]
max_features =["auto", "sqrt", "log2"]
#grid
grid=dict(n_estimators=n_estimators,criterion=criterion,max_features=max_features)
#cv
from sklearn.model_selection import RepeatedStratifiedKFold
cv=RepeatedStratifiedKFold(n_splits=5,n_repeats=3,random_state=1)
#GridSearchCV
from sklearn.model_selection import GridSearchCV
grid_cv=GridSearchCV(estimator=model,param_grid=grid,cv=cv,scoring='accuracy')
#results
res=grid_cv.fit(xtrain,ytrain)
print("best parameters are :",res.best_params_)
print("best accuracy is :",res.best_score_)
```

best parameters are : {'criterion': 'entropy', 'max_features': 'auto', 'n_estimators': 10}
best accuracy is : 0.128125

In []:

In []:

In []:

In []:

In []: