

In [ ]:

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
```

```
In [4]: df=pd.read_excel("Company_Data.csv.xlsx")
```

```
In [6]: df.head()
```

```
Out[6]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	Education	Urban
0	9.50	138.0	73.0	11.0	276.0	120.0	Bad	42.0	17.0	Yes
1	11.22	111.0	48.0	16.0	260.0	83.0	Good	65.0	10.0	Yes
2	10.06	113.0	35.0	10.0	269.0	80.0	Medium	59.0	12.0	Yes
3	7.40	117.0	100.0	4.0	466.0	97.0	Medium	55.0	14.0	Yes
4	4.15	141.0	64.0	3.0	340.0	128.0	Bad	38.0	13.0	Yes

```
In [23]: df.drop(['ShelveLoc', 'Urban', 'US'], axis=1, inplace=True)
```

```
In [24]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sales           400 non-null    float64
1   CompPrice       400 non-null    float64
2   Income          400 non-null    float64
3   Advertising     400 non-null    float64
4   Population      400 non-null    float64
5   Price           400 non-null    float64
6   Age             400 non-null    float64
7   Education       400 non-null    float64
dtypes: float64(8)
memory usage: 25.1 KB
```

```
In [25]: df.isna().sum()
```

```
Out[25]: Sales           0
CompPrice           0
Income             0
Advertising         0
Population          0
Price              0
Age                0
Education           0
dtype: int64
```


```
In [36]: x=df.drop(columns=["Education"])
y=df["Education"]
```

```
In [37]: from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.2,random_state=0)
```

```
In [38]: #Bagging Meta Estimator Classifier
from sklearn.ensemble import BaggingClassifier
#making the bagging classifier with 100 decision trees
model=BaggingClassifier(n_estimators=100)
#fitting data to bagging model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))
```

accuracy is : 0.1375

	precision	recall	f1-score	support
10.0	0.36	0.31	0.33	13
11.0	0.14	0.12	0.13	8
12.0	0.09	0.09	0.09	11
13.0	0.14	0.11	0.12	9
14.0	0.18	0.33	0.24	6
15.0	0.00	0.00	0.00	8
16.0	0.11	0.11	0.11	9
17.0	0.06	0.33	0.11	3
18.0	0.00	0.00	0.00	13
accuracy			0.14	80
macro avg	0.12	0.16	0.13	80
weighted avg	0.13	0.14	0.13	80



## Implementing The Bagging Classifier with Logistic Model

```
In [39]: #Bagging Meta Estimator Classifier
from sklearn.ensemble import BaggingClassifier
from sklearn.linear_model import LogisticRegression
#making the bagging classifier with 100 Logistic Regression models
model=BaggingClassifier(base_estimator=LogisticRegression(),n_estimators=100)
#fitting data to bagging model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))
```

C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
C:\Users\Admin\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

## Random Forest Classifier

```
In [40]: from sklearn.ensemble import RandomForestClassifier
model=RandomForestClassifier(n_estimators=100)
#fitting data to Random Forest model
model.fit(xtrain,ytrain)
#testing on test dataset
ypred=model.predict(xtest)
#Model Evaluation
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
print("accuracy is :",accuracy_score(ytest,ypred))
cm=confusion_matrix(ytest,ypred)
sns.heatmap(cm,annot=True)
print(classification_report(ytest,ypred))
```

```
accuracy is : 0.1875
              precision    recall  f1-score   support

   10.0         0.33         0.31         0.32         13
   11.0         0.33         0.38         0.35          8
   12.0         0.08         0.09         0.09         11
   13.0         0.12         0.11         0.12          9
   14.0         0.18         0.33         0.24          6
   15.0         0.00         0.00         0.00          8
   16.0         0.22         0.22         0.22          9
   17.0         0.13         0.67         0.22          3
   18.0         0.00         0.00         0.00         13

 accuracy          0.19         0.19         0.19         80
 macro avg         0.16         0.23         0.17         80
 weighted avg         0.16         0.19         0.16         80
```



## Tuning HyperParameters of Random Forest

```
In [42]: #model
model=RandomForestClassifier()
n_estimators =[10,50,100,1000]
criterion =["gini", "entropy"]
max_features =["auto", "sqrt", "log2"]
#grid
grid=dict(n_estimators=n_estimators,criterion=criterion,max_features=max_features)
#cv
from sklearn.model_selection import RepeatedStratifiedKFold
cv=RepeatedStratifiedKFold(n_splits=5,n_repeats=3,random_state=1)
#GridSearchCV
from sklearn.model_selection import GridSearchCV
grid_cv=GridSearchCV(estimator=model,param_grid=grid,cv=cv,scoring='accuracy')
#results
res=grid_cv.fit(xtrain,ytrain)
print("best parameters are :",res.best_params_)
print("best accuracy is :",res.best_score_)
```

best parameters are : {'criterion': 'entropy', 'max\_features': 'auto', 'n\_estimators': 10}  
best accuracy is : 0.128125

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: